

Xiao Yan

Email: Xiao.Yan@utdallas.edu

Phone: 214-549-0732

Location: Dallas, TX

Personal Website

Department of Computer Science, UT Dallas

Education

University of Texas at Dallas	08/2024 – Present
<i>Ph.D. in Computer Science, Systems Track — Mobile Computing, On-Device Intelligence, Edge-Cloud Systems</i>	
Imperial College London	10/2022 – 10/2023
<i>M.Sc. Applied Computational Science & Engineering with Merit</i>	
University of Birmingham	08/2019 – 07/2022
<i>B.Sc. Computer Science with Honours, First Class</i>	

Publications

Published

- **X. Yan** et al., “Are We There Yet? A Measurement Study of Efficiency for LLM Applications on Mobile Devices,” in *Proc. 2nd International Workshop on Foundation Models for Cyber-Physical Systems & IoT (FMSys '25)*, 2025.

Under Review

- **X. Yan** et al., “Hestia: Semantic Acoustic Sensing and On-Device LLM Reasoning for Privacy-Preserving Home Routine Recognition” (under review, ACM MobiCom).

Research Experience

Multimodal Edge Sensing for Wildlife Monitoring (Field Deployment)	06/2026 – Present
---	-------------------

The University of Texas at Dallas

- Leading the multimodal systems track of a wildlife-monitoring field deployment combining camera, acoustic, PIR, and satellite-RF (GNSS) sensing on autonomous solar-powered edge nodes.
- Prototyped a two-tier architecture moving computation toward the sensing front-end: sub-watt ESP32 sensor nodes gate wake-ups of a Raspberry Pi 5 running MobileNetV2 inference, cutting projected energy several-fold versus always-on capture.
- Built and bench-validated the full pipeline end-to-end, using instrumented power profiling of each sensor and compute stage to guide the hardware/software co-design.

Privacy-Preserving Home Routine Recognition via Acoustic Sensing	08/2025 – 02/2026
---	-------------------

The University of Texas at Dallas

- Designed Hestia, a hybrid edge-central system that reconstructs daily routines from sparse acoustic detections using commodity microphones and on-device LLM reasoning, with no raw audio leaving the home.
- Developed hierarchical temporal clustering, bounded on-device LLM inference within a deterministic scaffold, and continual self-adaptation via LLM pseudo-labels to improve edge models over time.
- Deployed and evaluated across two real homes, achieving routine recognition accuracy comparable to or exceeding camera-, WiFi-, and IMU-based baselines while preserving user privacy.

Efficiency of LLM Applications on Mobile Devices	01/2025 – 03/2025
---	-------------------

The University of Texas at Dallas

- Evaluated efficiency tradeoffs across mobile, edge, and cloud deployments for LLM-based applications, benchmarking latency, energy, and accuracy under multiple model compression strategies.
- Built AutoLife-Lite, a lightweight LLM application that infers user location and activity contexts from smartphone sensor data in real time.

Emotion-Aware Chatbot with Facial Expression Recognition	09/2021 – 03/2022
---	-------------------

University of Birmingham

- Built a multimodal chatbot fusing real-time CNN facial expression recognition with LSTM dialogue modeling; deployed end-to-end as a standalone application.

Technical Skills

Programming Languages: Python, Java, C/C++, MATLAB, JavaScript, HTML/CSS

Frameworks & Libraries: PyTorch, TensorFlow, NumPy, Pandas, Scikit-learn

Tools & Platforms: Git, Linux, LaTeX, Docker, Android Studio, Raspberry Pi, ESP32

Domains: Mobile Computing, On-Device LLM Systems, Edge-Cloud Architecture, Acoustic Sensing

Languages: English (Professional), Mandarin Chinese (Native)